

EECE 5550 Mobile Robotics HW #2

Instructor: Michael Everett

Due: Feb 7, 2023 at 11:00am

This assignment should be submitted via Gradescope. Your answers to both problems 1 & 2 should be submitted as the written portion. We do not plan to run your code this time (i.e., no autograder), but please also upload your code to Gradescope anyway.

Question 1: Bayesian inference with linear-Gaussian models

In this exercise we will study Bayesian estimation in linear-Gaussian models; as we will see later in the course, these play a fundamental role in robotic state estimation (most prominently in the celebrated [Kalman filter](#)).

We begin by recording a few useful facts. Recall that:

$$X \sim \mathcal{N}(\mu, \Sigma) \tag{1}$$

means that $X \in \mathbb{R}^n$ is a random variable that follows a Gaussian distribution with mean $\mu \in \mathbb{R}^n$ and covariance $\Sigma \in \mathbb{S}_{++}^n$. As we saw in class, X is described by the following probability density function:

$$p_X : \mathbb{R}^n \rightarrow \mathbb{R}$$
$$p_X(x) \triangleq \frac{1}{\sqrt{\det(2\pi\Sigma)}} \exp\left(-\frac{1}{2}(x - \mu)^\top \Sigma^{-1}(x - \mu)\right). \tag{2}$$

If we expand the quadratic form and ignore the normalization constant in (2), we find:

$$p_X(x) \propto \exp\left(-\frac{1}{2}\left(x^\top \Sigma^{-1}x - 2\mu^\top \Sigma^{-1}x\right)\right). \tag{3}$$

It follows from (3) that *any* function of the form:

$$f(z) = \frac{1}{c} \exp\left(-\frac{1}{2}\left(z^\top \Lambda z - 2\eta^\top z\right)\right) \tag{4}$$

with $c > 0$ is an unnormalized density for a Gaussian random variable $Z \sim \mathcal{N}(\bar{\mu}, \bar{\Sigma})$ with parameters:

$$\bar{\Sigma} = \Lambda^{-1}, \quad \bar{\mu} = \bar{\Sigma}\eta. \tag{5}$$

Equations (4) and (5) give an alternative way of parameterizing a Gaussian probability density, called the *information* or *canonical form*.

Now, suppose that Θ is a random variable with prior distribution:

$$\Theta \sim \mathcal{N}(\mu_0, \Sigma_0), \tag{6}$$

and that we collect a set of m noisy linear measurements $\tilde{Y}_1, \dots, \tilde{Y}_m$ of Θ according to:

$$\tilde{Y}_i = A_i \Theta + b_i + \epsilon_i, \quad \epsilon_i \sim \mathcal{N}(\mu_i, \Sigma_i), \quad (7)$$

where A_i , b_i , μ_i , and Σ_i are known parameters for all $i = 1, \dots, m$. In this exercise, you will determine the posterior distribution for Θ given the measurements $\tilde{Y}_1, \dots, \tilde{Y}_m$.

- (a) Use Bayes' Rule to express the posterior density $p(\Theta | \tilde{Y}_1, \dots, \tilde{Y}_m)$ in terms of the prior $p(\Theta)$ and the measurement likelihoods $p(\tilde{Y}_i | \Theta)$ for each individual measurement. You may leave your result in an unnormalized form.
- (b) Derive an expression for the likelihood function $p(\tilde{Y}_i | \Theta)$ of the i th measurement. (Hint: Notice that you can easily solve (7) for ϵ_i .)
- (c) Using your results from parts (a) and (b), derive the parametric form of the posterior density $p(\Theta | \tilde{Y}_1, \dots, \tilde{Y}_m)$. You should simplify your result by collecting linear and quadratic terms in Θ in the exponent. You may leave your result in an unnormalized form.
(Hint: Since your result need not be normalized, any term appearing in an exponent that does *not* involve Θ can be discarded by absorbing it into the normalization constant. You can use this fact to dramatically simplify your work.)
- (d) You should be able to recognize your expression for $p(\Theta | \tilde{Y}_1, \dots, \tilde{Y}_m)$ in part (c) as an unnormalized Gaussian density in information form. This shows that the posterior distribution for Θ is Gaussian; that is, $\Theta | \tilde{Y}_1, \dots, \tilde{Y}_m \sim \mathcal{N}(\bar{\mu}, \bar{\Sigma})$ for some mean $\bar{\mu}$ and covariance $\bar{\Sigma}$. What are the mean $\bar{\mu}$ and covariance $\bar{\Sigma}$ of this distribution?

Problem 2: Route planning in occupancy grid maps

As we will see later in the semester, occupancy grid maps provide a convenient representation of a robot's environment that is particularly well-suited to route planning for navigation.

For example, Fig 1a shows a (probabilistic) occupancy grid map of a research lab constructed using the [Cartographer](#) SLAM system¹, and Fig. 1b the resulting estimate of free and occupied space obtained by thresholding these occupancy probabilities to binary values.

In this exercise, you will implement two of the graph-based planning algorithms that we discussed in class (A* search and probabilistic road maps) to perform route-planning in the (binary) occupancy grid map shown in Fig. 1b.

Note: Since the A* search algorithm requires the use of several data structures other than basic matrices (e.g. sets and priority queues), we recommend implementing the following exercise in a Python notebook (feel free to use colab to manage your dependencies again!).

- (a) **A* search:** In the first part of this exercise, you will implement a general version of A* search that is abstracted with respect to the choice of representation of the graph G . This will enable you to apply it to *both* occupancy grids (considered as 8-connected graphs) *and* “standard” graphs G (for use with probabilistic roadmaps).


The pseudocode for this version of A* search is shown as Algorithm 1. The following objects appear in Algorithm 1:

- “CostTo” is a [map](#) that assigns to each vertex v the cost of the shortest known path from the start node s to v .

¹We will see later in the course how to construct these maps from raw sensor data.



(a) Probabilistic occupancy grid map



(b) Binary occupancy grid map

- “pred” is a map that associates to each vertex v its predecessor on the shortest known path from the start s to v .
- “EstTotalCost” is a map that assigns to each vertex v the sum $\text{CostTo}(v) + h(v, g)$, the sum of the cost of the best known path to v and the predicted cost of the best path from v to the goal g ; this is the estimated cost of the optimal path from the start s to the goal g that passes *through* vertex v .
- Q is a priority queue in which elements with *lower* values are removed *first*.
- “RecoverPath” is a function that takes as input the start state s , goal state g , and populated predecessor map pred, and returns the sequence of vertices on the optimal path from s to g .

Tip: In Python, you can use the [dictionary](#) class to implement the maps CostTo, pred, and EstTotalCost, a [list](#) or [set](#) to hold the vertex set V , and a sorted list or [heap](#) of (priority, vertex) tuples² to implement the priority queue Q .

- Implement the RecoverPath function.
 - Using your implementation of RecoverPath, implement the complete A* search algorithm shown in Algorithm 1. Your code should accept as input the vertex set V , the start and goal vertices s and g , and function handles for N , w , and h .
- (b) **Route planning in occupancy grids with A* search:** In this part of the problem, you will apply your A* search algorithm to perform route planning directly on the occupancy grid map Fig. 1b, considered as an 8-connected graph. Note that in this part of the exercise, we will *identify* each vertex $v \in V$ (cell) using its row and column (r, c) in the occupancy grid.
- Given an occupancy grid map M in the form of a 2D binary array, where a value of 0 indicates “occupied” and a value of 1 indicates “free” space, implement the function $N(v)$ that returns the set of unoccupied neighbors of a vertex v (remember that we can’t drive the robot through occupied space!) The vertex v and its neighbors should be expressed in the form of (row, column) tuples $v = (r, c)$.

²Python compares tuples in lexicographic order, so placing the priority first ensures that (priority, vertex) tuples will be sorted by priority, as desired.

Algorithm 1 An abstracted implementation of A^* search

Input: Graph $G = (V, E)$ with vertex set V and edge set E , nonnegative weight function $w: V \times V \rightarrow \mathbb{R}_+$ for the edges, admissible A^* heuristic $h: V \times V \rightarrow \mathbb{R}_+$ that returns the estimated cost-to-go, a set-valued function $N: V \rightarrow 2^V$ that returns the neighbors of a vertex v in G , starting vertex $s \in V$, goal vertex $g \in V$.

Output: A least-cost path from s to g if one exists, or the empty set \emptyset if no path exists.

```
1: function A_STAR_SEARCH( $V, s, g, N, w, h$ )
    // Initialization
2:   for  $v \in V$  do
3:     Set CostTo[ $v$ ] =  $+\infty$ .
4:     Set EstTotalCost[ $v$ ] =  $+\infty$ .
5:   end for
6:   Set CostTo[ $s$ ] = 0.                                ▷ Cost to reach starting vertex  $s$  is 0.
7:   Set EstTotalCost[ $s$ ] =  $h(s, g)$ .                  ▷ Estimated cost-to-go from  $s$  to  $g$ 
8:   Initialize  $Q = \{(s, h(s, g))\}$ .                  ▷ Insert start vertex  $s$  with value  $h(s, g)$ 
    // Main loop
9:   while  $Q$  is not empty do
10:     $v = Q.\text{pop}()$                                 ▷ Remove least-value element from  $Q$ 
11:    if  $v = g$  then                                ▷ We have reached the goal!
12:      return RECOVERPATH( $s, g, \text{pred}$ )              ▷ Reconstruct and return optimal path
13:    end if
14:    for  $i \in N(v)$  do                                ▷ For each of  $v$ 's neighbors
15:       $\text{pvi} = \text{CostTo}[v] + w(v, i)$                   ▷ Cost of path to reach  $i$  through  $v$ 
16:      if  $\text{pvi} < \text{CostTo}[i]$  then
        // The path to  $i$  through  $v$  is better than the previously-known best path to  $i$ ,
        // so record it as the new best path to  $i$ .
17:      Update  $\text{pred}[i] = v$ 
18:      Update  $\text{CostTo}[i] = \text{pvi}$                       ▷ Update cost of best path to  $i$ 
19:      Update  $\text{EstTotalCost}[i] = \text{pvi} + h(i, g)$ 
20:      if  $Q$  contains  $i$  then
21:         $Q.\text{setPriority}(i) = \text{EstTotalCost}[i]$       ▷ Update  $i$ 's priority
22:      else
23:         $Q.\text{insert}(i, \text{EstTotalCost}[i])$           ▷ Insert  $i$  into  $Q$  with priority  $\text{EstTotalCost}[i]$ 
24:      end if
25:    end if
26:  end for
27: end while
28: return  $\emptyset$                                     ▷ Return empty set: there is no path to goal
29: end function
```

- (ii) We will consider the cost of moving from a cell $v_1 = (r_1, c_1)$ to an adjacent cell $v_2 = (r_2, c_2)$ to be the Euclidean distance between the cell centers. Implement a function $d: V \times V \rightarrow \mathbb{R}_+$ that accepts as input the tuples v_1 and v_2 , and returns this Euclidean distance.
- (iii) We saw in class that the straight-line Euclidean distance between two points provides an admissible A^* heuristic h for route-planning using the total path length as the cost; this means that we can do route planning using your distance function d from part (ii) as both the edge weight w and the heuristic h .

Using your implementations of d , N , and A^* search, find the shortest path in the occupancy grid Fig. 1b from the starting point $s = (635, 140)$ to the goal $g = (350, 400)$ [assuming 0-based indexing for rows and columns, as is standard in CS.] Plot this optimal path overlaid on the image, and calculate its total length.

Tip: In Python, you can use the [Python Imaging Library](#) to easily manipulate basic image data. The following code snippet will read the occupancy map file from disk, interpret it as a Python [numpy](#) array, and then threshold it to produce the binary array M required in part (i):

- (c) **Route planning with probabilistic roadmaps:** Voxelized grids (like occupancy maps) provide simple and convenient models of robot configuration spaces, but as we will see in class their memory requirements scale *exponentially* in the dimension of the state space, making them far too costly to use for higher-dimensional planning problems.

Thus, *sampling-based planners* can sometimes provide a tractable alternative for planning in high-dimensional spaces. Recall that these methods *approximate* the configuration space C using a graph $G = (V, E)$ whose vertex set $V \subset C$ is a *randomly sampled subset* of points in C , and where two vertices $v_1, v_2 \in V$ are joined by an edge if v_2 is *reachable* from v_1 by applying a local controller.

In this part of the exercise, you will implement a sampling-based planner to perform route planning in the occupancy grid shown in Fig. 1b; more specifically, you will implement a *probabilistic roadmap* (PRM). Recall that we construct a PRM incrementally by sampling a new vertex $v_{new} \in C$, and then attempting to join v_{new} to nearby vertices $v \in G$ using a local planner (cf. Algorithms 2 and 3). In order to implement this approach, we must therefore specify:

- A method for sampling new vertices $v_{new} \in C$ (line 4 of Alg. 2)
- A suitable distance function $d: V \times V \rightarrow \mathbb{R}_+$ for characterizing “nearby” vertices (line 3 of Alg. 3)
- A local planner (line 4 of Alg. 3)

Algorithm 2 Construction of a probabilistic roadmap

Input: Desired number of sample points N , maximum local search radius d_{max} .

Output: A graph $G = (V, E)$ consisting of a vertex set $V \subseteq C$ of cardinality N , and edge set E indicating reachability via local control.

```

1: function CONSTRUCTPRM( $N, d_{max}$ )
2:   Initialize  $V = \emptyset, E = \emptyset$ .
3:   for  $k = 1, \dots, N$  do
4:     Sample a new vertex  $v_{new} \in C$ .
5:     ADDVERTEX( $G, v_{new}, d_{max}$ )
6:   end for
7:   return  $G = (V, E)$ 
8: end function
```

- (i) Implement a function that accepts as input the occupancy grid map M , and returns a vertex $v = (r, c)$ sampled *uniformly randomly* from the free space in M . [Hint: consider rejection sampling with a uniform proposal distribution.]

Algorithm 3 Adding a vertex v_{new} to the probabilistic roadmap $G = (V, E)$

```
1: function ADDVERTEX( $G, v_{new}, d_{max}$ )
2:    $V \leftarrow V \cup \{v_{new}\}$ . ▷ Add vertex  $v_{new}$  to  $G$ 
3:   for  $v \in V$  satisfying  $v \neq v_{new}$  and  $d(v, v_{new}) \leq d_{max}$  do ▷ Link  $v_{new}$  to nearby vertices
4:     Attempt to plan a path from  $v_{new}$  to  $v$ .
5:     if planning succeeds then
6:        $E \leftarrow E \cup \{(v, v_{new})\}$  ▷ Add edge  $e = (v, v_{new})$  to  $G$ 
7:     end if
8:   end for
9: end function
```

- (ii) We saw in class that it is easy to plan straight-line paths between arbitrary points using a differential drive robot, since the robot can rotate in-place to face the correct direction before beginning to move. Therefore, we might consider using a *straight-line path planner* as our local planner in line 4 of Alg. 3. Using this approach, a point $v_2 \in V$ is reachable from v_1 if and only if the line segment joining v_1 and v_2 does not intersect any occupied cells in M .

Implement a function that performs this reachability check. Your function should accept as input the occupancy grid map M and two grid cells $v_1 = (r_1, c_1)$ and $v_2 = (r_2, c_2)$, and return a Boolean value indicating whether the line segment joining v_1 and v_2 in M is obstacle-free.

- (iii) With the aid of your results from parts (c)(i) and (c)(ii), and your distance function implementation from (b)(iii), implement Algorithm 2 in the form of a function that accepts as input an occupancy grid map M , the desired number of samples N , and the maximum local search radius d_{max} , and returns a PRM G constructed from M .

Tip: You may find it convenient to model the PRM G using the [Graph](#) class in Python's [NetworkX](#) library. If you do so, you can record the location (row and column) of each vertex v by setting its `pos` attribute. Similarly, given any straight-line path joining two vertices $v_1, v_2 \in V$ found in line 4 of Alg. 3, you can store the length of this path as the *weight* of the edge $e_{12} = (v_1, v_2)$ joining v_1 and v_2 in G . The following code snippet provides a minimal working example:

- (iv) Using your implementation of Algorithm 2, construct a PRM on the occupancy grid in Fig. 1b with $N = 2500$ samples and a maximum local search radius of $d_{max} = 75$ voxels. Plot the resulting graph overlaid on Fig. 1b. [Hint: you may find NetworkX's [draw_networkx](#) function useful here.]
- (v) Recall that given a PRM G for a configuration space C , and start $s \in C$ and goal $g \in C$, we can plan a route from s to g by first *adding* s and g to the PRM, and then searching for a shortest path from s to g in G .

Using the PRM you constructed in part (v), find a path from $s = (635, 140)$ to $g = (350, 400)$. [Note: If s and g initially lie in separate connected components of G , you may need to sample and add more vertices to G until s and g are path-connected.] Plot this path overlaid on Fig. 1b, and calculate its total length.

Tip: In part (v), you may use NetworkX's implementation of [A* search](#).